

CS-GY 6613 - Artificial Intelligence INET
Assignment 2 Report

1. Data Sources and Prep

- a. Corpus video: YouTube ID YcvECxtXoxQ (Toyota RAV4 review). Downloaded as input_video.mp4.
- b. Frame sampling: ffmpeg -vf "fps=1/5" → 1 frame every 5 seconds. Produced 559 frames (frame_0001.jpg ... frame_0559.jpg) covering 0–2,790 s. Stored in assignments/assignment-2/frames/.
- c. Query set: Hugging Face dataset aegean-ai/rav4-exterior-images (65 exterior frames at 5 s spacing, stored remotely; loaded on the fly in Part C). Each row includes image, timestamp, and metadata.

2. Detector Choice & Training (Part A) – part_a_finetune_seg_model.py

- a. Model: Ultralytics YOLO v26 segmentation (yolo26n-seg.pt) fine-tuned on the Carparts segmentation dataset.
 - i. Rationale: Pretrained part-level labels, fast inference, and segmentation masks (though downstream we keep boxes).
- b. Training configs:
 - i. Local notebook-friendly run: 1 epochs, 512 img size, batch 8, workers 4, AMP off on Apple M-series (device = "mps" fallback to CPU) (Used AI help to get memory safe defaults, see the script for details)
 - ii. GPU run for final model:
 - 1. 20 and 50 epochs runs on Lightning.AI A100 (script variant in comments). Resulting checkpoints saved as part_a_best.pt and part_a_best_50_epochs.pt respectively. Both copied locally for inference and uploaded to Github as well (Stable artifacts reside in assignments/assignment-2/models/)

3. Detection Corpus (Part B) – part_b_retrieval_corpus.py

- a. Inputs: sampled frames + part_a_best_50_epochs.pt.
- b. Inference settings: confidence threshold 0.25; uses YOLO segmentation but stores bounding boxes; records frame index and timestamp (derived from 5 s stride).
- c. Output schema (per detection): video_id, frame_index, timestamp_sec, class_id, class_label, confidence_score, bounding_box ([x_min, y_min, x_max, y_max]), individual box coords, frame_file, detector_name.
 - i. See README for description of schema
- d. Main artifact: video_detections_50_epochs.parquet

- i. 1,100 rows; 370 unique timestamps; frame range 1–559; time range 0–2,790 s.
- ii. Top classes (counts): trunk 179, wheel 170, hood 117, front_glass 88, right_mirror 76, front_bumper 71, front_left_light 60, back_glass 56, back_bumper 55, back_left_light 53.
- iii. Confidence: mean 0.47, median 0.40.
- iv. Intervalization (5 s gap rule) by class (count / avg duration in seconds): trunk 79 / 4.7; wheel 41 / 8.3; hood 47 / 5.7; front_glass 23 / 13.3; front_bumper 21 / 11.9; right_mirror 23 / 8.7; others mostly short (many single-frame, 0–2 s).
- e. Secondary artifact (20-epoch model): video_detections.parquet with 911 rows; kept for comparison but not used downstream.

4. Retrieval Logic (Part C) – `part_c_semantic_search.py`

- a. Steps:
 - i. Load part_a_best_50_epochs.pt.
 - ii. Read corpus detections from video_detections_50_epochs.parquet.
 - iii. Build a class→interval index: for each class, unique timestamps are merged into contiguous intervals if the gap ≤ 5 s (frame stride). Supporting detection count tallied per interval.
 - iv. For each query image (65 total), run detector (threshold 0.25) to obtain labels.
 - v. For each detected label, return all matching intervals from the index; if no matching intervals exist, emit a null interval; if no labels are detected, emit a fully null record.
- b. Output: query_retrieval_results.parquet
 - i. Rows: 7,550 (one row per query×interval).
 - ii. Coverage: 65 queries; 4 queries produced no intervals (indices 30, 36, 38, 58).
 - iii. Interval rows: 7,546 with timestamps; average interval duration 5.48 s (median 0.0 s because many single-frame hits), max 130 s; supporting detections mean 2.58, max 32.
 - iv. Label usage (rows): wheel 1,599; trunk 948; back_bumper 920; hood 846; back_glass 490; front_glass 460; front_bumper 399; others tail off.
 - v. Intervals per query: mean 124, median 122, max 235 — reflects the label-only matching strategy.
- c. Conclusion is retrieval is class-label exact-match with no visual embeddings or spatial reasoning and temporal grouping follows frame stride.

5. Steps to reproduce

- a. (Optional - I uploaded the model weights I used to GitHub) Re-train detector using part_a_finetune_seg_model.py

- i. For full 50-epoch quality, run the commented Lightning variant on a CUDA box and place the resulting part_a_best_50_epochs.pt into assignments/assignment-2/models/.
- b. Build detection corpus
 - i. Download & sample video as in README and ensure frames/ exists
 - ii. Run part_b_retrieval_corpus.py to build detection corpus
- c. Run semantic search
 - i. Run part_c_semantic_search.py to create a car part (class_label) to intervals from detection corpus mapping and use that to run semantic search

6. Failure cases and limitations

- a. Temporal resolution limited to 5 s sampling; brief appearances can be missed or collapse to 0 s intervals
- b. Retrieval is label-only: no visual similarity check within the class; results can include every interval containing that label, leading to many intervals per query and potential false positives.
- c. Some classes underrepresented (doors, tailgate) and some intervals are single-frame, giving zero-duration windows.
- d. Confidence threshold fixed at 0.25; no calibration or per-class tuning; masks unused (only boxes preserved).
- e. Dataset/domain gap: query set and corpus both from the same video family but still show slight distribution shift; four queries produced no detections.